

Homework 5: Recurrent Neural Networks

Student Name	Student ID
米家龙	18342075

- Homework 5: Recurrent Neural Networks
 - Exercise 1: Backpropagation through Time
 - ex.1 (a)
 - ex.1 (b)
 - Exercise 2: Vanishing/Exploding Gradients in RNNs
 - ex.2 (a)
 - ex.2 (b)
 - ex.2 (c)
 - Exercise 3: LSTMs
 - ex.3 (a)
 - ex.3 (b)
 - ex.3 (c)

Exercise 1: Backpropagation through Time

Consider the RNN (Recurrent Neural Network) as:




Figure 1: A recurrent neural network.

Each state h_t is given by:

$$h_t = \sigma(Wh_{t-1} + Ux_t), \text{ where } \sigma(z) = \frac{1}{1 + \exp(-z)}$$

Let L be a loss function defined as the sum over the losses L_t at every time step until time T :

$$L = \sum_{t=0}^T L_t, \text{ where } L_t \text{ is a scalar loss depending on } h_t.$$

In the following, we want to derive the gradient of this loss function with respect to the parameter W .

ex.1 (a)

Suppose we have $y = \sigma(Wx)$ where $y \in \mathbb{R}^n$, $x \in \mathbb{R}^d$ and $W \in \mathbb{R}^{n \times d}$. Derive the Jacobian

$$\frac{\partial y}{\partial x} = \text{diag}(\sigma') W \in \mathbb{R}^{n \times d}$$

根据已知条件，可以通过链式法则得到：

$$\begin{aligned} \frac{\partial y}{\partial x} &= \frac{\partial \sigma(Wx)}{\partial x} \\ &= \frac{\partial \sigma(Wx)}{\partial Wx} \frac{\partial Wx}{\partial x} \end{aligned} \quad (1)$$

对于 Wx 有：

$$Wx = \begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1m} \\ w_{21} & w_{22} & \cdots & w_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{31} & w_{32} & \cdots & w_{3m} \\ w_{n1} & w_{n2} & \cdots & w_{nm} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} = \begin{bmatrix} w_{11}x_1 + w_{12}x_2 + \cdots + w_{1m}x_m \\ w_{21}x_1 + w_{22}x_2 + \cdots + w_{2m}x_m \\ \vdots \\ w_{n1}x_1 + w_{n2}x_2 + \cdots + w_{nm}x_m \end{bmatrix} \quad (2)$$

假设 $p_i = w_{i1}x_1 + w_{i2}x_2 + \cdots + w_{im}x_m$ 因此对于 $\frac{\partial Wx}{\partial x}$ ：

$$\begin{aligned} \frac{\partial Wx}{\partial x} &= \begin{bmatrix} \frac{\partial p_1}{\partial x_1} & \frac{\partial p_1}{\partial x_2} & \cdots & \frac{\partial p_1}{\partial x_m} \\ \frac{\partial p_2}{\partial x_1} & \frac{\partial p_2}{\partial x_2} & \cdots & \frac{\partial p_2}{\partial x_m} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial p_n}{\partial x_1} & \frac{\partial p_n}{\partial x_2} & \cdots & \frac{\partial p_n}{\partial x_m} \end{bmatrix} \\ &= \begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1m} \\ w_{21} & w_{22} & \cdots & w_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{31} & w_{32} & \cdots & w_{3m} \\ w_{n1} & w_{n2} & \cdots & w_{nm} \end{bmatrix} = W \end{aligned} \quad (3)$$

对于 $\frac{\partial \sigma(Wx)}{\partial Wx}$ ，令 $y_i = \sigma(p_i)$ ， $i = 1, 2, \dots, n$ 则有：

$$\begin{aligned}
\frac{\partial \sigma(Wx)}{\partial Wx} &= \frac{\partial y}{\partial p} \\
&= \begin{bmatrix} \frac{\partial y_1}{p_1} & \frac{\partial y_2}{p_1} & \dots & \frac{\partial y_n}{p_1} \\ \frac{\partial y_2}{p_1} & \frac{\partial y_2}{p_2} & \dots & \frac{\partial y_2}{p_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_n}{p_1} & \frac{\partial y_n}{p_2} & \dots & \frac{\partial y_n}{p_n} \end{bmatrix} \\
&= \begin{bmatrix} \sigma'(p_1) & 0 & \dots & 0 \\ 0 & \sigma'(p_2) & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \sigma'(p_n) \end{bmatrix} \\
&= \text{diag}(\sigma')
\end{aligned} \tag{4}$$

根据式(1), (3), (4)可得

$$\frac{\partial y}{\partial x} = \text{diag}(\sigma')W \in \mathbb{R}^{n \times d}$$

ex.1 (b)

Derive the quantity $\frac{\partial L}{\partial W} = \sum_{t=0}^T \sum_{k=1}^t \frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$.

根据 $L = \sum_{t=0}^T L_t$ 的定义, 可以得到

$$\begin{aligned}
\frac{\partial L}{\partial W} &= \frac{\partial}{\partial W} \sum_{t=0}^T L_t \\
&= \sum_{t=0}^T \frac{\partial L_t}{\partial W}
\end{aligned} \tag{5}$$

由于 L_t 是和 h_t 有关的一个标量损失, 并且 h_t 有如下表示:

$$\begin{cases} h_t &= \sigma(Wh_{t-1} + Ux_t) \\ h_{t-1} &= \sigma(Wh_{t-2} + Ux_{t-1}) \\ \vdots & \\ h_1 &= \sigma(Wh_0 + Ux_1) \\ h_0 &= \sigma(Wh_f + Ux_0), \text{ } h_f \text{ 和 } W \text{ 无关} \end{cases} \tag{6}$$

因此可以得到:

$$\begin{aligned}
\frac{\partial L_t}{\partial W} &= \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial W} \right) + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial W} \right) + \dots + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t-1}} \dots \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial W} \right) + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t-1}} \dots \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial W} \right) \\
&= \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_t} \frac{\partial h_t}{\partial W} \right) + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial W} \right) + \dots + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_1} \frac{\partial h_1}{\partial W} \right) + \left(\frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_0} \frac{\partial h_0}{\partial W} \right) \\
&= \sum_{k=0}^t \frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}
\end{aligned} \tag{7}$$

根据式 (5), (7) 可以得到:

$$\frac{\partial L}{\partial W} = \sum_{t=0}^T \sum_{k=0}^t \frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W} \tag{8}$$

发现结果和题目有矛盾, k 应该从0而不是1开始

Exercise 2: Vanishing/Exploding Gradients in RNNs

In this exercise, we want to understand why RNNs (Recurrent Neural Networks) are specially prone to the Vanishing/Exploding Gradients problem and what role the eigenvalues of the weight matrix play. Consider part (b) of exercise 1 again.

ex.2 (a)

Write down $\frac{\partial L}{\partial W}$ as expanded sum for $T = 3$. You should see that if we want to backpropagate through n timesteps, we have to multiply the matrix $\text{diag}(\sigma')W$ n times with itself.

由式 (7) 可得, 当 $T = 3$ 时, 有:

$$\begin{aligned}\frac{\partial L}{\partial W} &= \sum_{t=0}^{T=3} \sum_{k=0}^t \frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W} \\ &= \left(\frac{\partial L_0}{\partial h_0} \frac{\partial h_0}{\partial h_0} \frac{\partial h_0}{\partial W} \right) + \left(\frac{\partial L_1}{\partial h_1} \frac{\partial h_1}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_1}{\partial h_1} \frac{\partial h_1}{\partial h_1} \frac{\partial h_1}{\partial W} \right) \\ &\quad + \left(\frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial W} + \frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial h_2} \frac{\partial h_2}{\partial W} \right) \\ &\quad + \left(\frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_1} \frac{\partial h_1}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_3} \frac{\partial h_3}{\partial W} \right) \quad (9) \\ &= \left(\frac{\partial L_0}{\partial h_0} \frac{\partial h_0}{\partial W} \right) + \left(\frac{\partial L_1}{\partial h_1} \frac{\partial h_1}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_1}{\partial h_1} \frac{\partial h_1}{\partial W} \right) \\ &\quad + \left(\frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial W} + \frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial W} \right) \\ &\quad + \left(\frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_0} \frac{\partial h_0}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_1} \frac{\partial h_1}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial W} \right)\end{aligned}$$

根据式 (6) 可以得到:

$$\frac{\partial h_t}{\partial h_{t-1}} = \text{diag}(\sigma')W \in \mathbb{R}^{n \times d} \quad (10)$$

将上式合并, 可以得到:

$$\begin{aligned}\frac{\partial L}{\partial W} &= \left(\frac{\partial L_0}{\partial h_0} \frac{\partial h_0}{\partial W} \right) + \left[\frac{\partial L_1}{\partial h_1} (\text{diag}(\sigma')W) \frac{\partial h_0}{\partial W} + \frac{\partial L_1}{\partial h_1} \frac{\partial h_1}{\partial W} \right] \\ &\quad + \left[\frac{\partial L_2}{\partial h_2} (\text{diag}(\sigma')W)^2 \frac{\partial h_0}{\partial W} + \frac{\partial L_2}{\partial h_2} (\text{diag}(\sigma')W) \frac{\partial h_1}{\partial W} + \frac{\partial L_2}{\partial h_2} \frac{\partial h_2}{\partial W} \right] \quad (11) \\ &\quad + \left[\frac{\partial L_3}{\partial h_3} (\text{diag}(\sigma')W)^3 \frac{\partial h_0}{\partial W} + \frac{\partial L_3}{\partial h_3} (\text{diag}(\sigma')W)^2 \frac{\partial h_1}{\partial W} + \frac{\partial L_3}{\partial h_3} (\text{diag}(\sigma')W) \frac{\partial h_2}{\partial W} + \frac{\partial L_3}{\partial h_3} \frac{\partial h_3}{\partial W} \right]\end{aligned}$$

在式 (11) 看到, 当 $T = 3$ 时, 需要乘 $(\text{diag}(\sigma')W)$ 3次

ex.2 (b)

Remember that any diagonalizable (square) matrix M can be represented by its eigendecomposition $M = Q\Lambda Q^{-1}$ where Q is a matrix whose i -th column corresponds to the i -th eigenvector of M and Λ is a diagonal matrix with the corresponding eigenvalues placed on the diagonals. Recall that every eigenvector v_i satisfies this linear equation $Mv_i = \lambda_i v_i$, where $\lambda_i = \Lambda_{ii}$ is an eigenvalue of M . Proof by induction that for such a matrix the product $\prod_{i=1}^n M$ can be represented as: $M^n = Q\Lambda^n Q^{-1}$.

对于 $n = 1$ 时, 根据题目可得:

$$M^1 = Q\Lambda^1 Q^{-1} \quad (12)$$

对于 $n = t - 1$ 时, 假设:

$$M^{t-1} = Q\Lambda^{t-1} Q^{-1} \quad (13)$$

那么在 $n = t$ 时, 可以得到:

$$\begin{aligned} M^t &= M^{t-1} M \\ &= (Q\Lambda^{t-1} Q^{-1})(Q\Lambda Q^{-1}) \\ &= Q\Lambda^{t-1} \Lambda Q^{-1} \\ &= Q\Lambda^t Q^{-1} \end{aligned} \quad (14)$$

因此可以得到 $\prod_{i=1}^n M = Q\Lambda^n Q^{-1}$

ex.2 (c)

Consider the weight matrix $\begin{bmatrix} 0.58 & 0.24 \\ 0.24 & 0.72 \end{bmatrix}$. Its eigendecomposition is:

$$W = Q\Lambda Q^{-1} = \begin{bmatrix} -0.6 & -0.8 \\ -0.8 & 0.5 \end{bmatrix} \begin{bmatrix} 0.9 & 0 \\ 0 & 0.4 \end{bmatrix} \begin{bmatrix} -0.6 & -0.8 \\ -0.8 & 0.6 \end{bmatrix}$$

Calculate W^{30} . What do you observe? What happens in general if the absolute value of all eigenvalues of W is smaller than 1? What happens if the absolute value of any eigenvalue of W is larger than 1? What if all eigenvalues are 1?

根据式 (14) 可以得到:


$$\begin{aligned} W^{30} &= Q\Lambda^{30} Q^{-1} \\ &= \begin{bmatrix} -0.6 & -0.8 \\ -0.8 & 0.6 \end{bmatrix} \begin{bmatrix} 0.9^{30} & 0 \\ 0 & 0.4^{30} \end{bmatrix} \begin{bmatrix} -0.6 & -0.8 \\ -0.8 & 0.6 \end{bmatrix} \\ &= \begin{bmatrix} 0.015260 & 0.020348 \\ 0.020348 & 0.027130 \end{bmatrix} \end{aligned} \quad (15)$$

1. 从式 (15) 可以看出, Λ 得30次方会导致结果相对变小很多
2. 当 W 的所有特征值的绝对值都小于1时, Λ 对角线上的元素的绝对值也都会小于1, 此时对于 n 来说, 其递增会导致 Λ 指数递减, 最终导致 W^n 的值变得很小
3. 当 W 的所有特征值的绝对值都大于1时, Λ 对角线上的元素的绝对值也都会大于1, 此时对于 n 来说, 其递增会导致 Λ 指数递增, 最终导致 W^n 的值变得很大
4. 当 W 的所有特征值的绝对值都等于1时, Λ 对角线上的元素的绝对值也都会等于1, 此时对于 n 来说, 其无论怎么变化, Λ 不变, 最终 W^n 的值不变

Exercise 3: LSTMs

Recall the elements of a module in an LSTM and the corresponding computations, where \odot stands for pointwise multiplication. For a good explanation on LSTMs you can refer to

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>. Consider the LSTM in Figure 2.



$$\begin{aligned}
 f_t &= \sigma(W_f h_{t-1} + U_f x_t) \\
 i_t &= \sigma(W_i h_{t-1} + U_i x_t) \\
 o_t &= \sigma(W_o h_{t-1} + U_o x_t) \\
 \tilde{C}_t &= \tanh(W_c h_{t-1} + U_c x_t) \\
 C_t &= f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \\
 h_t &= o_t \odot \tanh(C_t)
 \end{aligned}$$

Figure 2: A Long Short Term Memory network.

ex.3 (a)

What do the gates f_t , i_t and o_t do?

- 遗忘门 f_t : 选择上一时刻状态 C_{t-1} 的部分信息融入当前状态 C_t , 选择性的将 C_{t-1} 中与当前词 x_t 无关的内容遗忘并保留有用的信息
- 输入门 i_t : 判断当前词 x_t 的信息, 选择性的进行保留进细胞当前状态 C_t 或者丢弃
- 输出门 o_t : 判断当前细胞状态 C_t 的有用信息, 并选择性的输出到隐层 h_t 中

ex.3 (b)

Which of the quantities next to the figure are always positive?

Let's now try to understand how this architecture approaches the vanishing gradients problem. To calculate the gradient $\frac{\partial L}{\partial \theta}$, where θ stands for the parameters (W_f, W_o, W_i, W_c) , we now have to consider the cell state C_t instead of h_t . Like h_t in normal RNNs, C_t will also depend on the previous cell states

C_{t-1}, \dots, C_0 , so we get a formula of the form:

$$\frac{\partial L}{\partial W} = \sum_{t=0}^T \sum_{k=1}^t \frac{\partial L_t}{\partial C_t} \frac{\partial C_t}{\partial C_k} \frac{\partial C_k}{\partial W}$$

where note that the real formula is a bit more complicated since C_t also depends on f_t , i_t and \tilde{C}_t , which in turn all depend on W , but this can be neglected.

对于 f_t , i_t 和 o_t 来说, 采用了 Sigmoid 函数, 而该函数值域为 $(0, 1)$, 因此符号总是为正的

对于 \tilde{C}_t 由于 Tanh 函数的值域为 $(-1, 1)$, 因此 \tilde{C}_t , h_t 和 C_t 都可能符号为负

ex.3 (c)

We know that $\frac{\partial C_t}{\partial C_k} = \prod_{i=k+1}^t \frac{\partial C_i}{\partial C_{i-1}}$. Let $f_t = 1$ and $i_t = 0$ such that $C_t = C_{t-1}$ for all t . What is the gradient $\frac{\partial C_t}{\partial C_k}$ in this case?

$$\frac{\partial C_t}{\partial C_k} = \prod_{i=k+1}^t \frac{\partial C_i}{\partial C_{i-1}} = \prod_{i=k+1}^t 1 = 1 \quad (16)$$

显然, 梯度为1